

# Full Information Maximum Likelihood

## Utrecht University Winter School: Missing Data in R



**Utrecht  
University**

Kyle M. Lang

Department of Methodology & Statistics  
Utrecht University

# Outline

---

Maximum Likelihood

Full Information Maximum Likelihood

Auxiliary Variables



# FIML Intuition

---

FIML is an ML estimation method that is robust to ignorable nonresponse.

- FIML partitions the missing information out of the likelihood function so that the model is only estimated from the observed parts of the data.

After a minor alteration to the likelihood function, FIML reduces to simple ML estimation.

- So, let's review ML estimation before moving forward.



# Maximum Likelihood Estimation

---

ML estimation simply finds the parameter values that are “most likely” to have given rise to the observed data.

- The *likelihood* function is just a probability density (or mass) function with the data treated as fixed and the parameters treated as random variables.
- Having such a framework allows us to ask: “Given that I’ve observed these data values, what parameter values most probably describe these data?”



# Maximum Likelihood Estimation

---

ML estimation is usually employed when there is no closed form solution for the parameters we seek.

- This is why you don't usually see ML used to fit general linear models.

After choosing a likelihood function, we iteratively optimize the function to produce the ML estimated parameters.

- In practice, we nearly always work with the natural logarithm of the likelihood function (i.e., the *loglikelihood*).



# ML Intuition

---

Let's say we have the following  $N = 10$  observations.

- We assume these data come from a normal distribution with a known variance of  $\sigma^2 = 1$ .
- We want to estimate the mean of this distribution,  $\mu$ .

```
(y <- rnorm(n = 10, mean = 5, sd = 1))
```

```
[1] 5.060983 3.364836 4.968344 6.696222 3.610013  
[6] 6.627266 4.165329 4.615346 4.537332 6.024850
```

# ML Intuition

---


In ML estimation, we would define different normal distributions.

- Every distribution would have  $\sigma^2 = 1$ .
- Each distribution would have a different value of  $\mu$ .

We then compare the observed data to those distributions and see which distribution best fits the data.



# ML Intuition






# Likelihoods

---

Suppose we have the following model:

$$Y \sim N(\mu, \sigma^2).$$



# Likelihoods

---

For a given  $Y_n$ , we have:

$$P(Y_n | \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(Y_n - \mu)^2}{2\sigma^2}}. \quad (1)$$

If we plug estimated parameters into Equation 1, we get the probability of observing  $Y_n$  given  $\hat{\mu}$  and  $\hat{\sigma}^2$ :

$$P(Y_n | \hat{\mu}, \hat{\sigma}^2) = \frac{1}{\sqrt{2\pi\hat{\sigma}^2}} e^{-\frac{(Y_n - \hat{\mu})^2}{2\hat{\sigma}^2}}. \quad (2)$$

Applying Equation 2 to all  $N$  observations and multiplying the results produces a *likelihood*:

$$\hat{L}(\hat{\mu}, \hat{\sigma}^2) = \prod_{n=1}^N P(Y_n | \hat{\mu}, \hat{\sigma}^2).$$



# Likelihoods

---

We generally want to work with the natural logarithm of Equation 2. Doing so gives the *loglikelihood*:

$$\begin{aligned}\hat{\mathcal{L}}(\hat{\mu}, \hat{\sigma}^2) &= \ln \prod_{n=1}^N P(Y_n | \hat{\mu}, \hat{\sigma}^2) \\ &= -\frac{N}{2} \ln 2\pi - N \ln \hat{\sigma} - \frac{1}{2\hat{\sigma}^2} \sum_{n=1}^N (Y_n - \hat{\mu})^2\end{aligned}$$

ML tries to find the values of  $\hat{\mu}$  and  $\hat{\sigma}^2$  that maximize  $\hat{\mathcal{L}}(\hat{\mu}, \hat{\sigma}^2)$ .

- Find the values of  $\hat{\mu}$  and  $\hat{\sigma}^2$  that are *most likely*, given the observed values of  $Y$ .

# Likelihoods

Suppose we have a linear regression model:

$$Y = \beta_0 + \beta_1 X + \varepsilon,$$

$$\varepsilon \sim N(0, \sigma^2).$$

This model can be equivalently written as:

$$Y \sim N(\beta_0 + \beta_1 X, \sigma^2)$$




Image retrieved from:

<http://www.seaturtle.org/mtn/archives/mtn122/mtn122p1.shtml>

# Likelihoods


---

For a given  $\{Y_n, X_n\}$ , we have:

$$P(Y_n|X_n, \beta_0, \beta_1, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(Y_n - \beta_0 - \beta_1 X_n)^2}{2\sigma^2}}. \quad (3)$$

If we plug our estimated parameters into Equation 3, we get the probability of observing  $Y_n$  given  $\hat{Y}_n = \hat{\beta}_0 + \hat{\beta}_1 X_n$  and  $\hat{\sigma}^2$ .

$$P(Y_n|X_n, \hat{\beta}_0, \hat{\beta}_1, \hat{\sigma}^2) = \frac{1}{\sqrt{2\pi\hat{\sigma}^2}} e^{-\frac{(Y_n - \hat{\beta}_0 - \hat{\beta}_1 X_n)^2}{2\hat{\sigma}^2}} \quad (4)$$



# Likelihoods

---

So, our final loglikelihood function would be the following:

$$\begin{aligned}\hat{\mathcal{L}}(\hat{\beta}_0, \hat{\beta}_1, \hat{\sigma}^2) &= \ln \prod_{n=1}^N P(Y_n | X_n, \hat{\beta}_0, \hat{\beta}_1, \hat{\sigma}^2) \\ &= -\frac{N}{2} \ln 2\pi - N \ln \hat{\sigma} - \frac{1}{2\hat{\sigma}^2} \sum_{n=1}^N (Y_n - \hat{\beta}_0 - \hat{\beta}_1 X_n)^2.\end{aligned}$$



# Example

---

```
## Fit a model:
out1 <- lm(ldl ~ bp + glu + bmi, data = diabetes)

## Extract the predicted values and estimated residual standard error:
yHat <- predict(out1)
s     <- summary(out1)$sigma

## Compute the row-wise probabilities:
pY <- dnorm(diabetes$ldl, mean = yHat, sd = s)

## Compute the loglikelihood, and compare to R's version:
sum(log(pY)); logLik(out1)[1]

[1] -2109.939
[1] -2109.93
```

# Multivariate Normal Distribution

---

The PDF for the multivariate normal distribution is:


$$P(\mathbf{Y}|\boldsymbol{\mu}, \Sigma) = \frac{1}{\sqrt{(2\pi)^P |\Sigma|}} e^{-\frac{1}{2}(\mathbf{Y}-\boldsymbol{\mu})^T \Sigma^{-1}(\mathbf{Y}-\boldsymbol{\mu})}.$$

So, the multivariate normal loglikelihood is:

$$\mathcal{L}(\boldsymbol{\mu}, \Sigma) = - \left[ \frac{P}{2} \ln(2\pi) + \frac{1}{2} \ln |\Sigma| + \frac{1}{2} \right] \sum_{n=1}^N (\mathbf{Y}_n - \boldsymbol{\mu})^T \Sigma^{-1} (\mathbf{Y}_n - \boldsymbol{\mu}).$$

Which can be further simplified if we multiply through by -2:

$$-2\mathcal{L}(\boldsymbol{\mu}, \Sigma) = [P \ln(2\pi) + \ln |\Sigma|] \sum_{n=1}^N (\mathbf{Y}_n - \boldsymbol{\mu})^T \Sigma^{-1} (\mathbf{Y}_n - \boldsymbol{\mu}).$$






# Steps of ML

---

1. Choose a probability distribution,  $f(Y|\theta)$ , to describe the distribution of the data,  $Y$ , given the parameters,  $\theta$ .
2. Choose some estimate of  $\theta$ ,  $\hat{\theta}^{(i)}$ .
3. Compute each row's contribution to the loglikelihood function by evaluating:  $\ln \left[ f \left( Y_n | \hat{\theta}^{(i)} \right) \right]$ .
4. Sum the individual loglikelihood contributions from Step 3 to find the loglikelihood value,  $\hat{\mathcal{L}}$ .
5. Choose a "better" estimate of the parameters,  $\hat{\theta}^{(i+1)}$ , and repeat Steps 3 and 4.
6. Repeat Steps 3 – 5 until the change between  $LL^{(i-1)}$  and  $LL^{(i)}$  falls below some trivially small threshold.
7. Take  $\hat{\theta}^{(i)}$  as your estimated parameters.



# FULL INFORMATION MAXIMUM LIKELIHOOD



# From ML to FIML

---

The  $n$ th observation's contribution to the multivariate normal loglikelihood function would be the following:

$$\mathcal{L}(\boldsymbol{\mu}, \boldsymbol{\Sigma})_n = -\frac{P}{2} \ln(2\pi) - \frac{1}{2} \ln |\boldsymbol{\Sigma}| - \frac{1}{2} (\mathbf{Y}_n - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{Y}_n - \boldsymbol{\mu}). \quad (5)$$



# From ML to FIML

---

The  $n$ th observation's contribution to the multivariate normal loglikelihood function would be the following:

$$\mathcal{L}(\boldsymbol{\mu}, \Sigma)_n = -\frac{P}{2} \ln(2\pi) - \frac{1}{2} \ln |\Sigma| - \frac{1}{2} (\mathbf{Y}_n - \boldsymbol{\mu})^T \Sigma^{-1} (\mathbf{Y}_n - \boldsymbol{\mu}). \quad (5)$$

FIML just tweaks Equation 5 a tiny bit:

$$\mathcal{L}(\boldsymbol{\mu}, \Sigma)_{fiml,n} = -\frac{P}{2} \ln(2\pi) - \frac{1}{2} \ln |\Sigma_q| - \frac{1}{2} (\mathbf{Y}_n - \boldsymbol{\mu}_q)^T \Sigma_q^{-1} (\mathbf{Y}_n - \boldsymbol{\mu}_q).$$


Where  $q = 1, 2, \dots, Q$  indexes response patterns.



# Visualize the Response Patterns

These data contain 4 unique response patterns.

- We'd define 4 different version of  $\mu$  and  $\Sigma$ .
- We'd calculate each individual loglikelihood contributions using the appropriate flavor of  $\mu$  and  $\Sigma$ .



# FIML Example

---

We most often apply FIML when estimating latent variable models.

- We can easily apply FIML to latent variable models in **lavaan**.

```
library(lavaan)

## Read in some data:
bfi <- readRDS(paste0(dataDir, "bfi_datasets.rds"))$incomplete

## Specify the measurement model:
cfaMod <- '
agree =~ A1 + A2 + A3 + A4 + A5
open  =~ O1 + O2 + O3 + O4 + O5
'

## Estimate the CFA using FIML:
fimlOut <- cfa(cfaMod, data = bfi, std.lv = TRUE, missing = "fiml")
```

# FIML Example

---

```
partSummary(fimlOut, 1:4)
```

lavaan 0.6.17 ended normally after 41 iterations

Estimator	ML
Optimization method	NLMINB
Number of model parameters	31
Number of observations	2800
Number of missing patterns	32

Model Test User Model:

Test statistic	360.865
Degrees of freedom	34
P-value (Chi-square)	0.000

# FIML Example

```
partSummary(fimlOut, 7, fmi = TRUE)
```

Latent Variables:

	Estimate	Std.Err	z-value	P(> z )	FMI
agree =~					
A1	0.538	0.038	14.109	0.000	0.399
A2	-0.764	0.030	-25.530	0.000	0.417
A3	-0.989	0.033	-29.722	0.000	0.434
A4	-0.693	0.039	-17.953	0.000	0.407
A5	-0.838	0.031	-26.622	0.000	0.403
open =~					
O1	0.635	0.025	24.968	0.000	0.003
O2	-0.640	0.036	-17.615	0.000	0.004
O3	0.831	0.029	28.880	0.000	0.008
O4	0.345	0.028	12.333	0.000	0.002
O5	-0.647	0.031	-20.890	0.000	0.005



# FIML Example

---

```
partSummary(fimlOut, 8, fmi = TRUE)
```

Covariances:

	Estimate	Std.Err	z-value	P(> z )	FMI
agree ~~					
open	-0.290	0.028	-10.530	0.000	0.147

# FIML Example

```
partSummary(fimlOut, 9, fmi = TRUE)
```

Intercepts:

	Estimate	Std.Err	z-value	P(> z )	FMI
.A1	2.422	0.031	77.128	0.000	0.281
.A2	4.814	0.025	189.726	0.000	0.252
.A3	4.616	0.028	163.443	0.000	0.232
.A4	4.735	0.032	146.063	0.000	0.275
.A5	4.573	0.027	167.561	0.000	0.243
.01	4.816	0.021	225.180	0.000	-0.000
.02	2.713	0.030	91.745	0.000	-0.000
.03	4.441	0.023	192.731	0.000	-0.000
.04	4.894	0.023	212.316	0.000	-0.000
.05	2.490	0.025	99.119	0.000	-0.000

# FIML Example

```
partSummary(fimlOut, 10, fmi = TRUE)
```

Variances:

	Estimate	Std.Err	z-value	P(> z )	FMI
.A1	1.693	0.059	28.649	0.000	0.352
.A2	0.765	0.037	20.491	0.000	0.463
.A3	0.736	0.047	15.632	0.000	0.489
.A4	1.654	0.061	27.307	0.000	0.354
.A5	0.876	0.043	20.256	0.000	0.484
.01	0.878	0.031	28.191	0.000	0.003
.02	2.039	0.062	32.972	0.000	0.005
.03	0.797	0.040	19.848	0.000	0.006
.04	1.369	0.038	35.732	0.000	0.002
.05	1.349	0.044	30.369	0.000	0.005
agree	1.000				
open	1.000				

# FMI with FIML

---

As you saw above, we can also estimate the FMI when using FIML.

- The FMI is calculated using the method described by Savalei and Rhemtulla (2012).

Savalei and Rhemtulla (2012) take an information-theoretic approach to defining the FMI.

- Based on the *Missing Information Principle* of Orchard and Woodbury (1972)
- Their FMI estimates the ratio of missing to complete information for each parameter.

You can use this method to compute the FMI for sufficient statistics via the **semTools::fmi()** function.



# AUXILIARY VARIABLES



# Satisfying the MAR Assumption

---


Like MI, FIML also requires MAR data.

- Parameters will be biased when data are MNAR.

Unlike MI, FIML directly treats the missing data while estimating the analysis model.

- The MAR predictors must be included in the analysis model.
- Otherwise, FIML reduces to pairwise deletion.

When the MAR predictors are not substantively interesting variables, naively including them in the analysis model can change the model's meaning.



# Saturated Correlates Technique


---

Graham (2003) developed the *saturated correlates* approach to meet two desiderata:

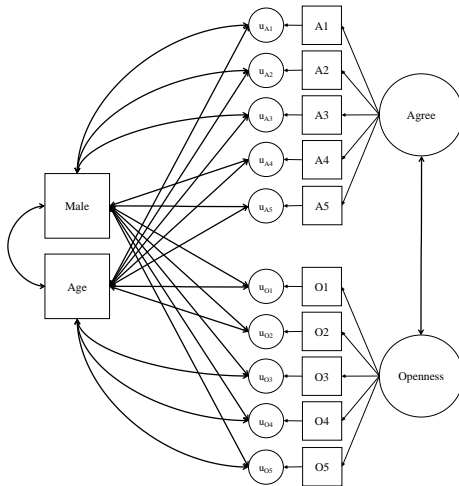
1. Satisfy the MAR assumption by incorporating MAR predictors into the analysis model.
2. Maintain the fit and substantive meaning of the analysis model.

The approach entails incorporating the MAR predictors via a fully-saturated covariance structure:

1. Allow every MAR predictor to covary with all other MAR predictors.
2. Allow every MAR predictor to covary with all observed variables in the analysis model (or their residuals).



# Saturated Correlates Diagram





# Saturated Correlates Example

---

We can use the `lavaan.auxiliary()` function from **semTools** (or one of its wrappers) to streamline the analysis.

```
library(semTools)

## Estimate the CFA from above with auxiliary variables:
fimlOut2 <- bfi %>%
  mutate(male = as.numeric(sex == "male")) %>%
  cfa.auxiliary(cfaMod,
               data = .,
               aux = c("age", "male"),
               std.lv = TRUE)
```

# Saturated Correlates Example

---

The `cfa.auxiliary()` function has automatically added the following paths to our CFA model.

```
age ~~ age
age ~~ male
age ~~ A1
age ~~ A2
age ~~ A3
age ~~ A4
age ~~ A5
age ~~ O1
age ~~ O2
age ~~ O3
age ~~ O4
age ~~ O5
```

```
male ~~ male
male ~~ A1
male ~~ A2
male ~~ A3
male ~~ A4
male ~~ A5
male ~~ O1
male ~~ O2
male ~~ O3
male ~~ O4
male ~~ O5
```

# Saturated Correlates Example

The auxiliaries have been correlated with all other variables.

```
inspect(fimlOut2, "est")$theta[11:12, 1:6] %>% round(3)
```

	A1	A2	A3	A4	A5	O1
age	-2.927	1.330	1.188	2.622	1.863	0.763
male	0.119	-0.094	-0.083	-0.082	-0.073	0.052

```
inspect(fimlOut2, "est")$theta[11:12, 7:12] %>% round(3)
```

	O2	O3	O4	O5	age	male
age	-0.773	0.491	0.147	-1.484	123.998	-0.246
male	-0.006	0.020	0.001	-0.015	-0.246	0.220

# Saturated Correlates Example

---

The degrees of freedom have not changed, though.

```
## Naive FIML:  
fitMeasures(fimlOut, "df")  
  
df  
34  
  
## FIML w/ saturated correlates:  
fitMeasures(fimlOut2, "df")  
  
df  
34
```

# Saturated Correlates Example

Let's compare the effects of the various missing data treatments on the latent covariance estimates.

	Complete Data	Listwise Deletion	Multiple Imputation	Naive FIML	FIML w/ Auxiliaries
Est	-0.306	-0.254	-0.299	-0.290	-0.295
FMI	—	—	0.305	0.147	0.146

Latent Covariances

# References

---

- Graham, J. W. (2003). Adding missing-data-relevant variables to fimo-based structural equation models. *Structural Equation Modeling: A Multidisciplinary Journal*, 10(1), 80–100. doi: 10.1207/S15328007SEM1001\_4
- Orchard, T., & Woodbury, M. A. (1972). A missing information principle: Theory and applications. In *Proceedings of the sixth berkeley symposium on mathematical statistics and probability, volume 1: Theory of statistics*.
- Savalei, V., & Rhemtulla, M. (2012). On obtaining estimates of the fraction of missing information from full information maximum likelihood. *Structural Equation Modeling: A Multidisciplinary Journal*, 19(3), 477–494.

